

CSIE 2136 Algorithm Design and Analysis, Fall 2020



Graph Algorithms - II

Hsu-Chun Hsiao

Announcement

- HW3 due in four weeks (12/24)
- Mini-hw8 due next week

3.5-week Agenda

- Graph basics
 - Graph terminology [B.4, B.5]
 - Real-world applications
 - Graph representations [Ch. 22.1]
- Graph traversal
 - Breadth-first search (BFS) [Ch. 22.2]
 - Depth-first search (DFS) [Ch. 22.3]
- DFS applications
 - Topological sort [Ch. 22.4]
 - Strongly-connected components [Ch. 22.5]
- Minimum spanning trees [Ch. 23]
 - Kruskal's algorithm
 - Prim's algorithm
- Single-source shortest paths [Ch. 24]
 - Dijkstra algorithm
 - Bellman-Ford algorithm
 - SSSP in DAG
- All-pairs shortest paths [Ch. 25]
 - Floyd-Warshall algorithm
 - Johnson's algorithm

Today's Agenda

- DFS applications
 - Topological sort [Ch. 22.4]
 - Strongly-connected components [Ch. 22.5]
- Minimum spanning trees [Ch. 23]
 - Kruskal's algorithm
 - Prim's algorithm
- Shortest paths: terminology and properties
 - Edge relaxation
 - Shortest-paths properties

Application of DFS: Topological Sort

Textbook chapter 22.4

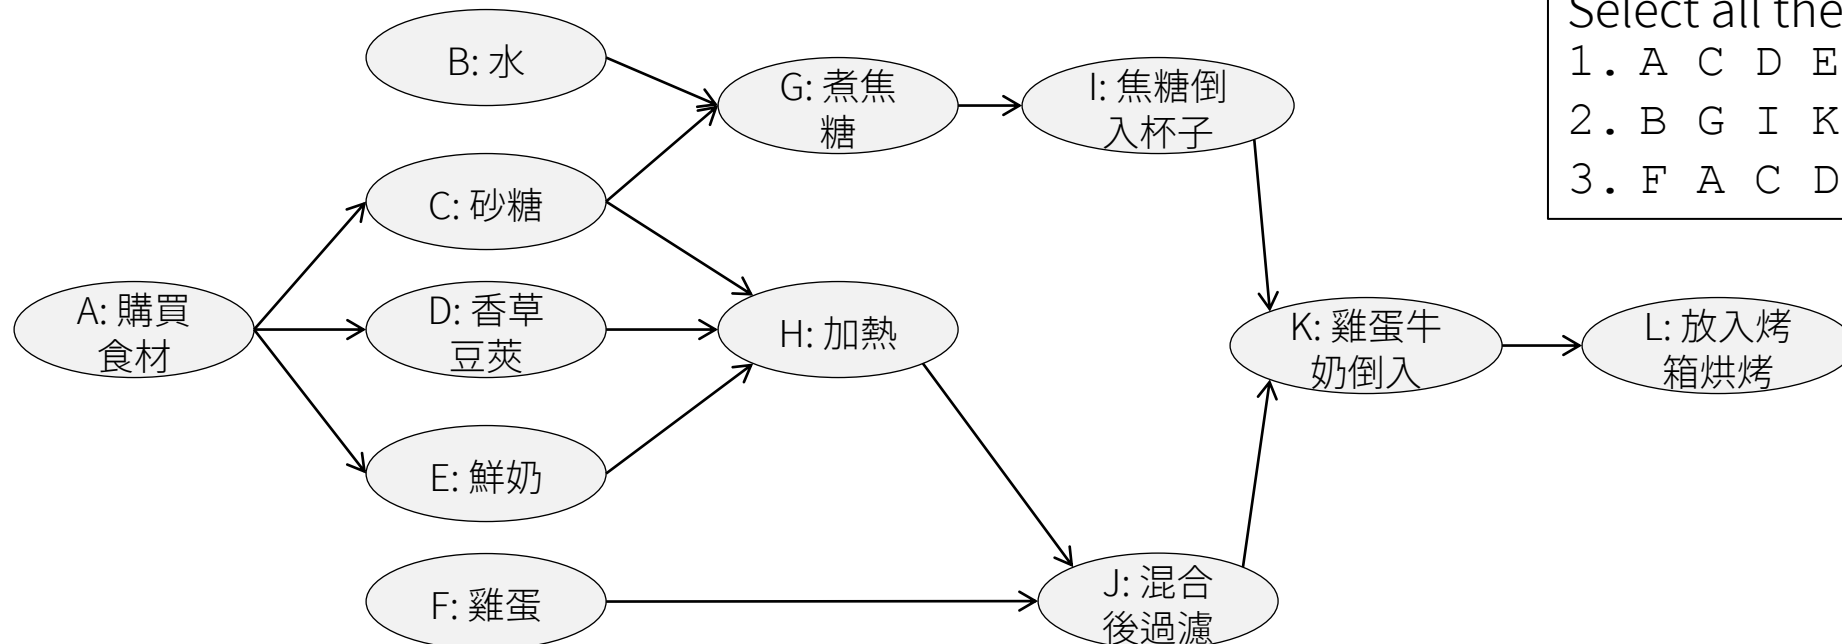


MasterChef: 布丁篇



Slido Poll:
#ADA2020

- A->B: 要先處理完A才能處理B
- 新手一次只能做一件事，用什麼順序才能順利做出布丁？
- Intuition: 前置作業要先完成，才能做後面的步驟

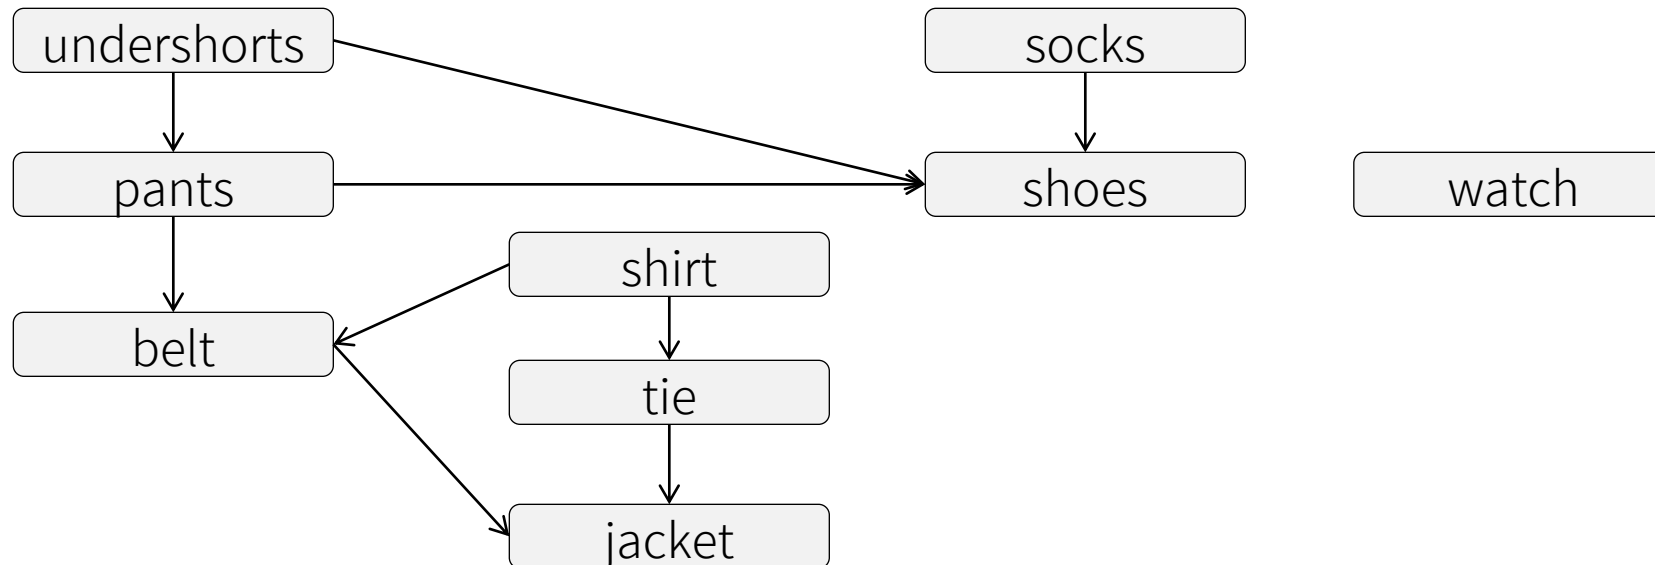


Select all the valid orders

1. A C D E H B G I F J K L
2. B G I K A C H E D F J L
3. F A C D E H B G I J K L


Directed Acyclic Graphs (DAGs)

- A **DAG** is a directed graph with no cycles
- Often used to indicate precedence among events (**X must happen before Y**)
 - E.g., cooking, taking courses, clothing...



Topological Sort

- Input: a DAG $G = (V, E)$
- Output: a linear ordering of all its vertices such that for all edges (u, v) in E , u precedes v in the ordering
- Alternative view: a vertex ordering along a horizontal line so that **all directed edges go from left to right**
- A DAG can have multiple valid topological orders
 - E.g., watch can be placed anywhere in the following example



Topological sort algorithm

```
TOPOLOGICAL-SORT(G) // G is a DAG
```

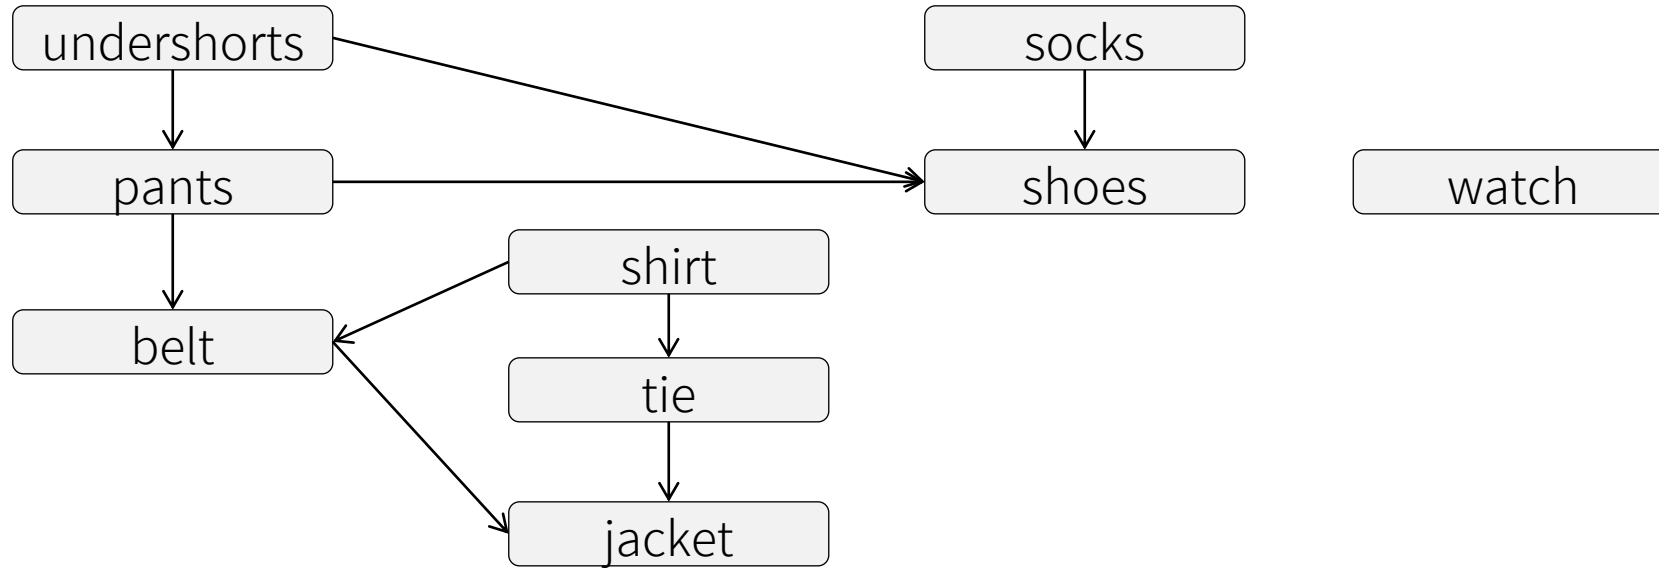
```
    Call DFS(G) to compute finishing times v.f for each vertex v
```

```
    As each vertex is finished, insert it onto the front of a linked list
```

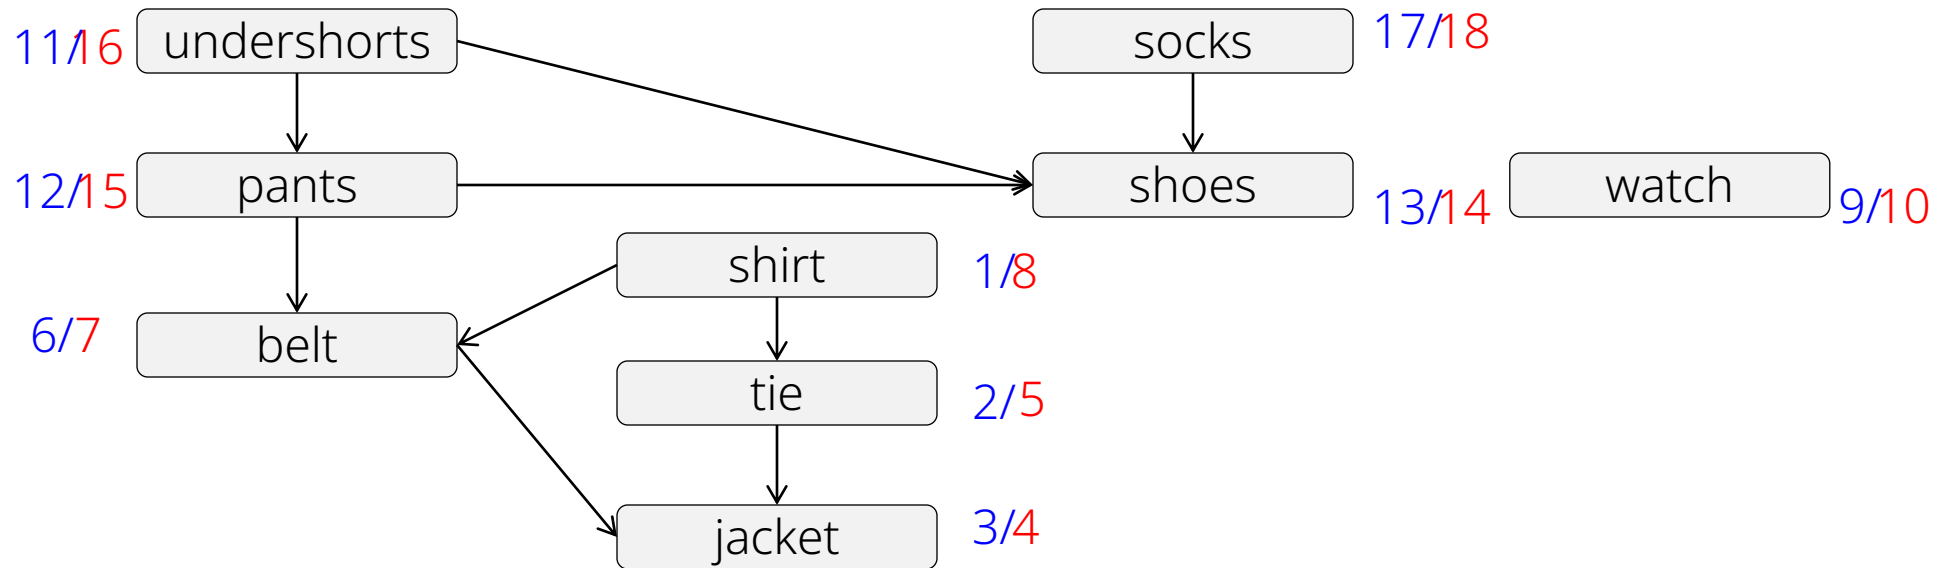
```
    return the linked list of vertices
```

- Vertex *u* is in front of *v* $\Leftrightarrow u.f > v.f$
- We will prove this linked list comprises a topological ordering

Topological sort using DFS



Topological sort using DFS



socks undershorts pants shoes watch shirt belt tie jacket

Running time analysis

```
TOPOLOGICAL-SORT(G) // G is a DAG
```

```
    Call DFS(G) to compute finishing times v.f for each vertex v
```

```
    As each vertex is finished, insert it onto the front of a linked list
```

```
    return the linked list of vertices
```

- DFS with adjacency lists: $\Theta(V + E)$ time
- Insert each vertex to the linked list: $\Theta(V)$ time
- \Rightarrow total running time is $\Theta(V + E)$

Another topological sort algorithm: Kahn's algorithm

- Intuition: removing “source vertices” one by one and updating in-degree values
 - Source vertices: vertices with in-degree = 0
- Running time is $\Theta(V + E)$
 - Need to maintain in-degree values and a queue of current source vertices

Lemma 22.11 Characterizing directed acyclic graphs

A directed graph is acyclic \Leftrightarrow a DFS yields no back edges

Proof by contradiction: the \Rightarrow direction

- Suppose there is a back edge (u, v)
 - $\Rightarrow v$ is an ancestor of u in DFS forest
 - \Rightarrow There is a path from v to u in G and (u, v) completes the cycle
 - \Rightarrow Contradiction!

Lemma 22.11 Characterizing directed acyclic graphs

A directed graph is acyclic \Leftrightarrow a DFS yields no back edges

Proof by contradiction (cont.): the \Leftarrow direction

- Suppose there is a cycle \mathcal{C}
 - \Rightarrow Let v be the first vertex in \mathcal{C} to be discovered and u is the predecessor of v in \mathcal{C}
 - \Rightarrow Upon discovering v the whole cycle from v to u is WHITE
 - \Rightarrow At time $v.d$, the vertices of \mathcal{C} form a path of WHITE vertices from v to u
 - \Rightarrow By the **white-path theorem**, vertex u becomes a descendant of v in the depth-first forest
 - \Rightarrow Therefore, (u, v) is a back edge
 - \Rightarrow Contradiction!

Theorem 22.12 Correctness of topological sort algorithm

The algorithm produces a topological sort of the input DAG

對所有的 edge (u, v) ，證明在此 list 中 u 一定在 v 前面（也就是 $u.f > v.f$ 成立）

Proof

- When (u, v) is explored, u is gray.
- Consider three cases of v : gray, white, black

Theorem 22.12 Correctness of topological sort algorithm


The algorithm produces a topological sort of the input DAG

Proof (cont.)

- $v = \text{gray}$
 - $\Rightarrow (u, v) = \text{back edge}$
 - $\Rightarrow G$ is cyclic (by Lemma 22.11)
 - \Rightarrow Contradiction, so v cannot be gray
- $v = \text{white}$
 - $\Rightarrow v$ becomes descendant of u (by white-path theorem)
 - $\Rightarrow v$ will be finished before u
 - $\Rightarrow v.f < u.f$
- $v = \text{black}$
 - $\Rightarrow v$ is already finished
 - $\Rightarrow v.f < u.f$

Cycle detection using DFS

- Since **cycle detection** becomes **back-edge detection** (Lemma 22.11), DFS can be used to test whether a graph is a DAG.





Slido Poll:
#ADA2020

Q: Is there a DFS forest for a cyclic graph?

Q: Is there a topological order for a cyclic graph?

Q: Given a topological order, is there always a DFS traversal that produces such an order?

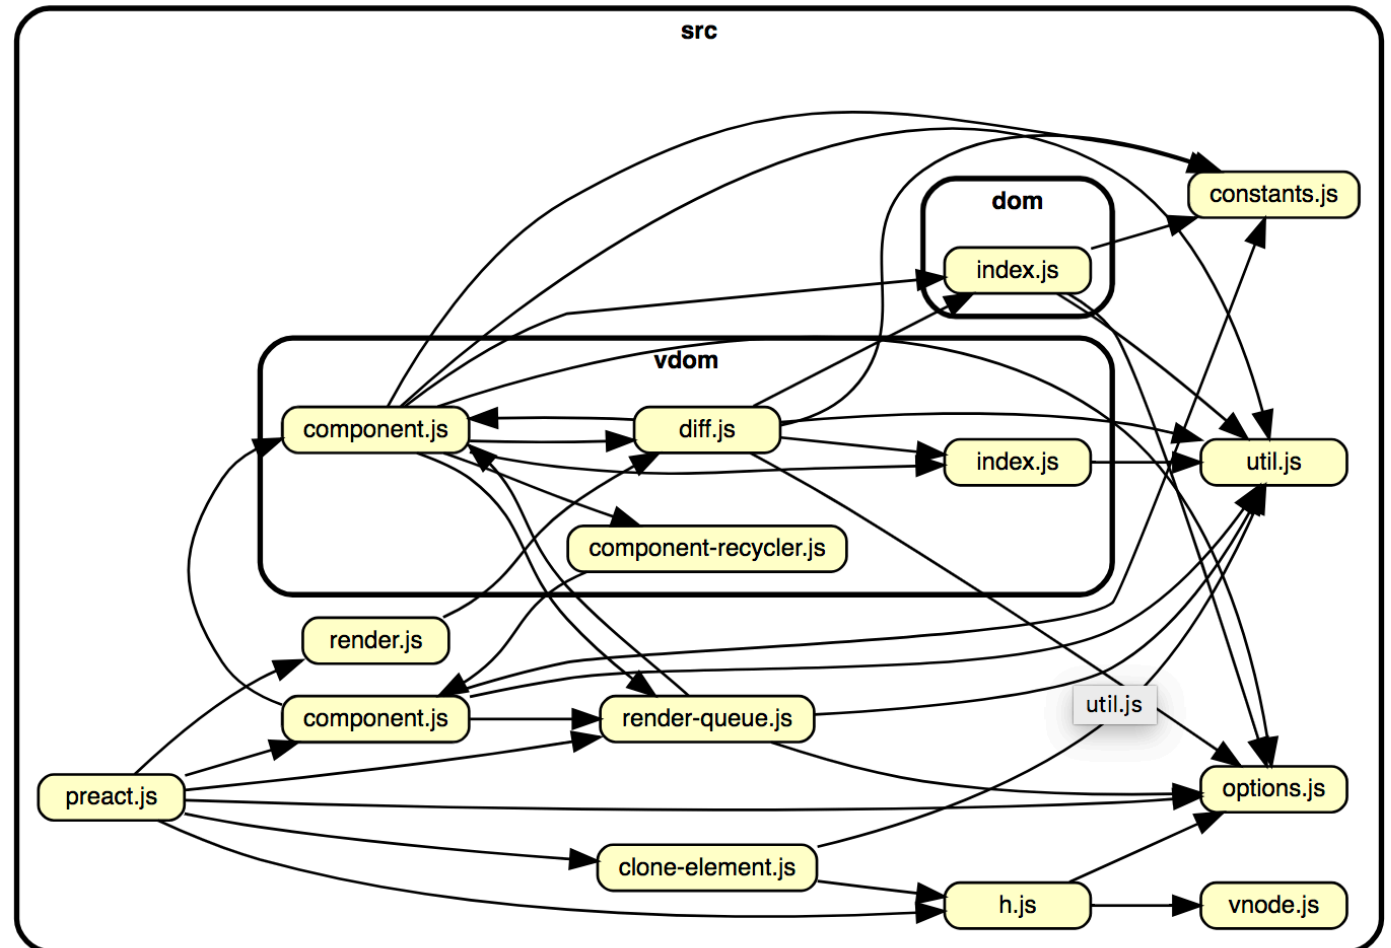
Strongly Connected Components (SCC)

Software module dependency graph

Vertex = software module
Edge = dependency

How to identify mutual dependency?


Another example of mutual dependency:
Brooklyn Nine-Nine - Amy Applies for a
Block-Party Request (Episode Highlight)
<https://youtu.be/FYM04gQAyr8>



<https://www.netlify.com/blog/2018/08/23/how-to-easily-visualize-a-projects-dependency-graph-with-dependency-cruiser/>

Connected components of an undirected graph

The connected components of an undirected graph are the equivalence classes of vertices under the “is reachable from” relation.




3 connected components: {A,B,E}, {C,F}, {D}

Strongly connected components of a directed graph

The strongly connected components of a directed graph are the equivalence classes of vertices under the “mutually reachable” relation.


That is, a strong component is a maximal subset of mutually reachable nodes.



4 strongly connected components: {A,B,E}, {C}, {D}, {F}


Decomposing a directed graph

- A directed graph is a DAG of its strongly connected components



$G = (V, E)$

Contract each SCC
into one vertex



Component graph $G^{scc} = (V^{scc}, E^{scc})$

Q: Show that a component graph must be a DAG

Q: Does the following algorithm determine whether a graph G is strongly connected in $O(V + E)$ time?



```
Run BFS in  $G$  from any node  $s$   
Run BFS in the transpose of  $G$ , from the same source node  $s$   
If both BFS executions found all nodes, return true; otherwise, return false
```

Yes

Note: we denote a **transpose** or **reverse** graph of a directed graph $G = (V, E)$ as G^T , and $G^T = (V, E^T)$ where $E^T = \{(v, u) \mid (u, v) \in E\}$

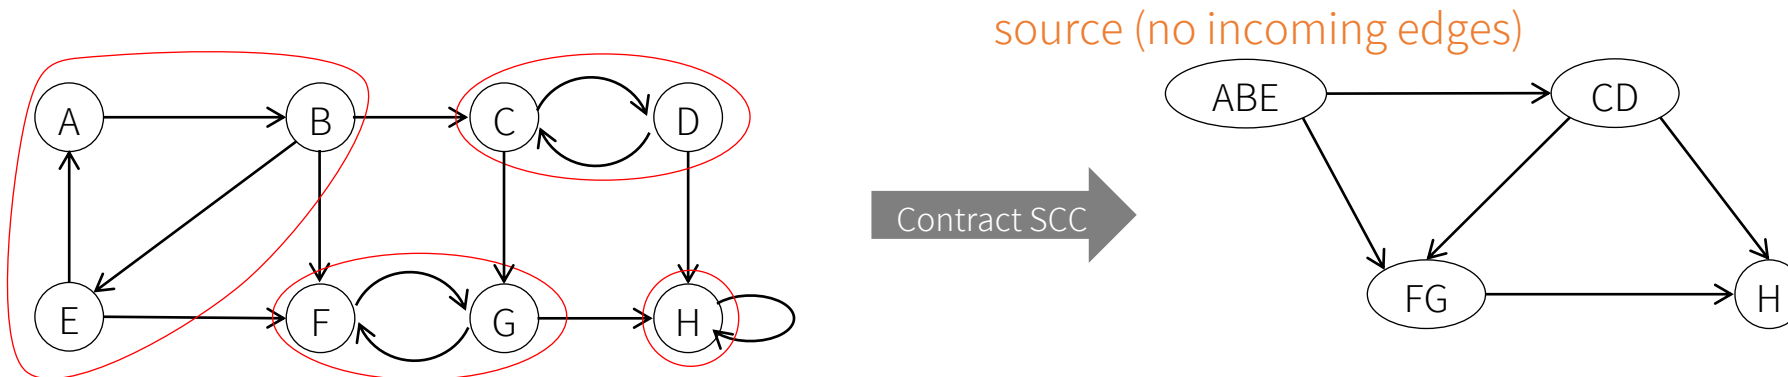
Finding SCC

- **Observation 1:** Starting from s , DFS finds all reachable nodes from s . Hence, if we can select a vertex in a **sink SCC** as the starting vertex for DFS, then DFS will discover **all (and only)** nodes in the sink SCC.
 - \Rightarrow we can find SCCs one by one in a reverse topological order of G^{scc} !
 - However, how to identify a vertex in a sink SCC?



Finding SCC

- **Observation 2 (Exercises 22.5-4):** An SCC in G is also an SCC in G^T . Also, a source SCC in G is a sink SCC in G^T .
- **Observation 3:** Finding a source SCC is easy. The vertex with the highest finishing time (found by running DFS in G) must be in a **source SCC**.
 - Implied by Lemma 22.14 (will prove it in a few slides)



Finding SCC: the Kosaraju-Sharir algorithm


`Strongly-Connected-Components(G)`

```
1  call  $DFS(G)$  to compute finishing times  $u.f$  for each vertex  $u$ 
2  compute  $G^T$ 
3  call  $DFS(G^T)$ , but in the main loop of DFS, consider the vertices in order of
   decreasing  $u.f$  (as computed in line 1)
4  output the vertices of each tree in the DFS forest formed in line 3 as a
   separate strongly connected component
```


- Time complexity
 - 2 DFS executions
 - $\Theta(V + E)$ using adjacency lists

Let's try it!

1 call $\text{DFS}(G)$ to compute u.f



2 compute G^T
3 call $\text{DFS}(G^T)$, in decreasing order of u.f



Lemma 22.14

Let C and C' be distinct strongly connected components in directed graph $G = (V, E)$. Suppose that there is an edge (u, v) where u in C and v in C' . Then $f(C) > f(C')$.

Here we define $f(U) = \max_{u \in U} \{u.f\}$, and $d(U) = \min_{u \in U} \{u.d\}$

Proof

- Consider two cases: $d(C) < d(C')$ and $d(C) > d(C')$
- If $d(C) < d(C')$:
 - Let x be the first vertex discovered in C
 - \Rightarrow At $t = x.d$, all vertices in C and C' are white
 - \Rightarrow At $t = x.d$, there is a white path from x to every vertex in C and C' (why?)
 - \Rightarrow By the white-path theorem, they are all x 's descendants in the DFS tree
 - \Rightarrow By the parenthesis theorem, $x.f$ is the largest
 - $\Rightarrow f(C) = x.f > f(C')$

Lemma 22.14

Let \mathcal{C} and \mathcal{C}' be distinct strongly connected components in directed graph $G = (V, E)$. Suppose that there is an edge (u, v) where u in \mathcal{C} and v in \mathcal{C}' .

Then $f(\mathcal{C}) > f(\mathcal{C}')$.

Here we define $f(U) = \max_{u \in U} \{u.f\}$, and $d(U) = \min_{u \in U} \{u.d\}$

Proof (cont'd)


- If $d(\mathcal{C}) > d(\mathcal{C}')$:
 - Let y be the first vertex discovered in \mathcal{C}'
 - \Rightarrow At $t = y.d$, all vertices in \mathcal{C}' are white
 - \Rightarrow At $t = y.d$, there is a white path from y to every vertex in \mathcal{C}'
 - \Rightarrow By the white-path theorem and the parenthesis theorem, all other vertices in \mathcal{C}' are y 's descendants and $y.f$ is the largest among them
 - $\Rightarrow f(\mathcal{C}') = y.f$
 - Because there is no path from \mathcal{C}' to \mathcal{C} (why?), no vertex in \mathcal{C} is reachable from y
 - \Rightarrow At $t = y.f$, all vertices in \mathcal{C} are still white
 - $\Rightarrow f(\mathcal{C}) > y.f > f(\mathcal{C}')$

Theorem 22.16 Correctness of the Kosaraju-Sharir algorithm


The Kosaraju-Sharir algorithm correctly computes the strongly connected components of the directed graph G provided as its input

Proof by induction on the number of DFS trees in line 3

- Inductive hypothesis: the first k trees produced are SCC
 - Base case: when $k = 0$, trivially correct
- Inductive step: assume the first k trees are SCC, consider the $(k + 1)$ th tree T
 - Let u be the first vertex of T , and let u be in SCC C
 - We will show that the vertices of T are the same as vertices in C



All vertices in C are in T :



All vertices in T are in C :

Theorem 22.16 Correctness of the Kosaraju-Sharir algorithm



The Kosaraju-Sharir algorithm correctly computes the strongly connected components of the directed graph G provided as its input

Proof by induction (cont'd)

- Inductive step: assume the first k trees are SCC, consider the $(k + 1)$ th tree T
 - Let u be the first vertex of T , and let u be in SCC C
 - We will show that the vertices of T are the same as vertices in C
 - All vertices in C are in T :

By the inductive hypothesis, at $t = u.d$, all other vertices of C are white.
By the white-path theorem, all vertices in C are descendants of u in T .
 - All vertices in T are in C :

By construction, $u.f$ is the largest among nodes that have yet to be visited in line 3.
That is, $u.f = f(C) > f(C')$, where C' is any SCC other than C that has yet to be visited.
Because there is no edge from C to C' in G^T (why?), T will not contain any vertices in any C' .





Slido Poll:
#ADA2020

Q: Can the following algorithms correctly find SCCs?

Strongly-Connected-Components-1(G)

- 1 **compute** G^T
- 2 call $DFS(G^T)$ to compute finishing times $u.f$ for each vertex u
- 3 call $DFS(G)$, but in the main loop of DFS, consider the vertices in order of decreasing $u.f$ (as computed in line 1)
- 4 output the vertices of each tree in the DFS forest formed in line 3 as a separate strongly connected component

Strongly-Connected-Components-2(G)



- 1 call $DFS(G)$ to compute finishing times $u.f$ for each vertex u
- 2 call $DFS(G)$, but in the main loop of DFS, consider the vertices in order of increasing $u.f$ (as computed in line 1)
- 3 output the vertices of each tree in the DFS forest formed in line 3 as a separate strongly connected component

Minimum Spanning Trees

Textbook Chapter 23

Spanning tree

- Spanning tree of a graph G = a subgraph that is a tree and connects all the vertices
 - Exactly $n - 1$ edges
 - Acyclic
- There can be many spanning trees of a graph







Slido Poll:
#ADA2020

Spanning tree


- **BFS** and **DFS** also generate spanning trees
 - BFS tree is typically “short and bushy”
 - DFS tree is typically “long and stringy”




Q: Can the above spanning trees be generated from BFS/DFS?

Minimum spanning tree (MST)

- A **minimum spanning tree** of a graph G is a spanning tree with **minimal weight**
- Weight of a tree T = the sum of weights of all edges in T



Weight = 18



Weight = 12, MST

Q: How to find a MST in an unweighted graph?

Any spanning tree is an MST in an unweighted graph

Q: Given a weighted graph G , can there be more than one MST?

Yes, consider an unweighted graph: every spanning tree is an MST.

But we will show that MST is unique if all edge weights are distinct.

Q: If the edge weights are all increased by the same constant, does an MST of the old graph remain an MST in the re-weighted graph?

Yes

Minimum spanning tree (MST)

- Finding an MST is an **optimization** problem
- Two **greedy** algorithms compute an MST:
 - **Kruskal's algorithm**: consider edges in **ascending order of weight**. At each step, select the next edge as long as it does not create cycle.
 - **Prim's algorithm**: start with any vertex s and **greedily grow a tree from s** . At each step, add the edge of the least weight to connect an isolated vertex.

Kruskal's algorithm


Kruskal(G)

start with $T = V$ (no edges)

for each edge in increasing order by weight

if adding edge to T does not create a cycle

then add edge to T



Weight = 12
MST

Running time depends on how the **cycle test** is implemented.
Using a **disjoint-set data structure**, running time = $O(E \log V)$
(will show the details later)

Prim's Algorithm


Prim(G)

Start with a tree T with one vertex (any vertex)

while T is not a spanning tree

 Find least-weight edge that connects T to a new vertex

 Add this edge to T



Weight = 12
MST

Running time depends on how finding least-weight edge implemented.
Using a binary min-heap, running time = $O(E \log V)$
(will show the details later)

MST uniqueness

Q: Given G , can there be more than one MST?

Yes

Q: Given G with distinct edge weights, can there be more than one MST?

No

MST Uniqueness

MST is unique if all edge weights are distinct

Proof by contradiction

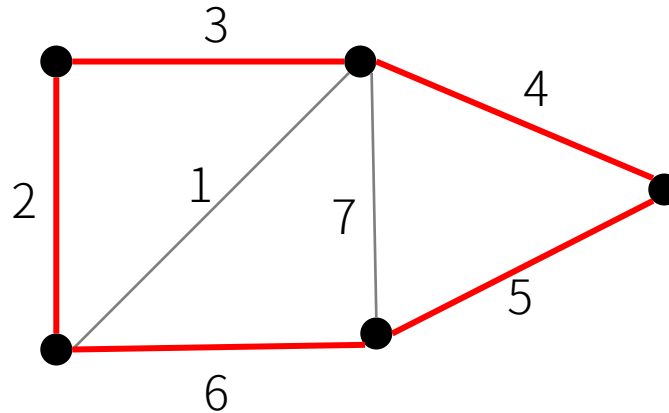
- Suppose there are two MSTs T_A and T_B on the same graph
 - Let e be the least-weight edge in $T_A \cup T_B$ and e is not in both
 - WLOG, assume e is in T_A
 - Add e to T_B
- $\Rightarrow \{e\} \cup T_B$ contains a cycle \mathcal{C}
- $\Rightarrow \mathcal{C}$ includes at least one edge e' that is not in T_A
- \Rightarrow In T_B , replacing e' with e yields a MST with less cost
- \Rightarrow Contradiction!

When edge costs are not distinct

- For proof purpose, we can break tie and ensure a unique MST by applying a **lexicographical order** of edges
- Define a **new weight function w'** over edges such that
 - $w'(e_i) < w'(e_j)$ if $w(e_i) < w(e_j)$ or $(w(e_i) = w(e_j)$ and $i < j)$
 - $w'(S_i) < w'(S_j)$ if $w(S_i) < w(S_j)$ or $(w(S_i) = w(S_j)$ and $S_i \setminus S_j$ has a lower indexed edge than $S_j \setminus S_i)$
- Hence, there is a unique MST w.r.t. to this new weight function w'
- **Note:** Prim and Kruskal algorithms don't require the weights to be distinct. The above is needed for the proof purpose only.

Cycle property

For simplicity, assume all edge weights are **distinct**, thus an unique MST. Let \mathcal{C} be any cycle in the graph G , and let e be an edge with the maximum weight on \mathcal{C} . Then **the MST does not contain e** .



No MST contains the edge of cost 6

Cycle property


For simplicity, assume all edge weights are **distinct**, thus an unique MST. Let \mathcal{C} be any cycle in the graph G , and let e be an edge with the maximum weight on \mathcal{C} . Then **the MST does not contain e** .

Proof by contradiction

- Suppose e is in the MST
- => Removing e disconnects the MST into two components T_1 and T_2
- => There exists another edge e' in \mathcal{C} that can reconnect T_1 & T_2
- => Since $\text{weight}(e') < \text{weight}(e)$, the new tree has a lower weight
- => Contradiction!

Cut property

For simplicity, assume all edge weights are **distinct**, thus an unique MST. Let \mathcal{C} be a cut (i.e., a partition of the vertices) in the graph, and let e be the edge with the minimum cost across \mathcal{C} . Then **the MST contains e** .



There is an MST containing the edge of cost 4

Cut property

For simplicity, assume all edge weights are **distinct**, thus an unique MST. Let \mathcal{C} be a cut (i.e., a partition of the vertices) in the graph, and let e be the edge with the minimum cost across \mathcal{C} . Then **the MST contains e** .

Proof by contradiction

- Suppose e is not in the current MST
- => Adding e creates a cycle in the MST
- => There exists another edge e' in the cut \mathcal{C} that can break the cycle
- => Since $\text{weight}(e') > \text{weight}(e)$, the new tree has a lower weight
- => Contradiction!

Kruskal's algorithm


`Kruskal(G)`

start with $T = V$ (no edges)

for each edge in increasing order by weight

if adding edge to T does not create a cycle

then add edge to T



Weight = 12
MST

Running time depends on how the **cycle test** is implemented.
Using a **disjoint-set data structure**, running time = $O(E \log V)$
(will show the details later)

Implementation of Kruskal's algorithm

```
MST-KRUSKAL(G,w) // w = weights
  A = empty // edge set of MST
  for v in G.V
    MAKE-SET(v)
  sort the edges of G.E into non-decreasing order by weight
  for (u,v) in G.E, taken in non-decreasing order by weight
    if FIND-SET(u) ≠ FIND-SET(v)
      A = A ∪ {u, v}
      UNION(u,v)
  return A
```

- Disjoint-set data structure: MAKE-SET, FIND-SET, UNION
- Each set contains the vertices in one tree of the current forest

Running time analysis

- The **amortized cost** of the disjoint-set-forest implementation with union-by-rank only (textbook Chapter 21):
 - MAKE-SET = $O(1)$
 - FIND-SET = $O(\log V)$
 - UNION = $O(\log V)$
 - The amortized cost of m operations on n elements is $O(m \log n)$ (Exercise 21.4-4)
- Sort edge = $O(E \log E) = O(E \log V)$
 - $\log E = O(\log V)$ because E is at most V^2
- Running time of Kruskal = $O(E \log V)$

Correctness of Kruskal's algorithm

Kruskal's algorithm computes the MST

Proof

- Consider whether adding e creates a cycle:
 1. If adding e to T creates a cycle \mathcal{C}
 - Then e is the max weight edge in \mathcal{C}
 - The **cycle property** ensures that e is not in the MST
 2. If adding $e = (u, v)$ to T does not create a cycle
 - Before adding e , the current set contains at least two trees T_1 and T_2 such that u in T_1 and v in T_2
 - e is the minimum cost edge on the cut of T_1 and T_2
 - The **cut property** ensures that e is in the MST

Prim's Algorithm


Prim(G)

Start with a tree T with one vertex (any vertex)

while T is not a spanning tree

 Find least-weight edge that connects T to a new vertex

 Add this edge to T



Weight = 12
MST

Running time depends on how finding least-weight edge implemented.
Using a binary min-heap, running time = $O(E \log V)$
(will show the details later)

Implementation of Prim's algorithm

```
MST-PRIM(G, w, r) // w = weights, r = root
  for u in G.V
    u.key =  $\infty$ 
    u. $\pi$  = NIL
  r.key = 0
  Q = G.V // BUILD-MIN-QUEUE
  while Q  $\neq$  empty
    u = EXTRACT-MIN(Q)
    for v in G.adj[u]
      if v  $\in$  Q and w(u, v) < v.key
        v. $\pi$  = u
        v.key = w(u, v) // DECREASE-KEY
```

- Q = min-priority queue, containing vertices not yet in the tree
- $v.key$ = minimum weight of any edge connecting v to the tree
- $v.\pi$ = the parent of v in the tree

Running time analysis

- Binary min-heap (textbook Chapter 6)
 - BUILD-MIN-HEAP = $O(V)$
 - EXTRACT-MIN = $O(\log V)$
 - DECREASE-KEY = $O(\log V)$
- Running time of Prim = $O(V \log V + E \log V)$
= $O(E \log V)$, because $V = O(E)$ in a connected graph
- Can be improved to $O(E + V \log V)$ using Fibonacci heaps (textbook Chapter 19)

Correctness of Prim's algorithm

Prim's algorithm computes the MST


Proof

1. Prove that all edges found by Prim's are in the MST:
 - Let S be the set of vertices in the current tree T
 - Prim's algorithm adds the cheapest edge e with exactly one endpoint in S
 - The **cut property** ensures that e is in the MST
2. Because Prim's outputs a spanning tree, $|\text{edges found by Prim's}| = n - 1$
 - \Rightarrow Edges found by Prim's = edges on the MST

Shortest Paths: Terminology and Properties


Textbook Chapter 24

訓練師在探索神獸級的神奇寶貝時被 Okapi 咬傷，每走一公尺就會損血一滴。請找出從金字塔 (s) 到荒野大夫家 (t) 的最短路徑？



Definitions


- Given a weighted, directed graph $G = (V, E)$
- Given a **weight function** w mapping an edge to a weight
 - Note that weights are arbitrary numbers, **not necessarily distances**
 - Weight function **needs not satisfy triangle inequality** (think about airline fares)
- Weight of path** $p = w(p) = \text{sum of weights of edges on } p$



The weight of path
 $s \rightarrow a \rightarrow c \rightarrow t$ is 23

Definitions

- **Shortest-path weight** $\delta(s, t)$ = minimum weight of path from s to t
- A **shortest path** from s to t = any path with weight $\delta(s, t)$





$\delta(s, t) = ?$

Shortest path from s to $t = ?$

Q: Can a shortest path contain a **negative-weight edge**?

Yes.



$\delta(s, v)$ remains well defined for all v , if G contains no negative-weight cycles reachable from the source s .



Q: Can a shortest path contain a **negative-weight cycle**?


Doesn't make sense.

If there is a negative-weight cycle on some path from s to v , we define $\delta(s, v) = -\infty$.



Q: Can a shortest path contain a positive-weight cycle?

No




$$p = \langle v_0, v_1, \dots, v_k \rangle$$
$$c = \langle v_i, v_{i+1}, \dots, v_j \rangle$$

Let $p' = \langle v_0, v_1, \dots, v_i, v_{j+1}, v_{j+2}, \dots, v_k \rangle$
 $w(p') \leq w(p)$ if $w(c) \geq 0$

Q: Can a shortest path contain a **zero-weight cycle**?

It may contain a zero-weight cycle, but then there must exist a simple path of the same weight.



$$p = \langle v_0, v_1, \dots, v_k \rangle$$
$$c = \langle v_i, v_{i+1}, \dots, v_j \rangle$$

Q: Can a shortest path contain a cycle?

We safely assume shortest paths have no cycles

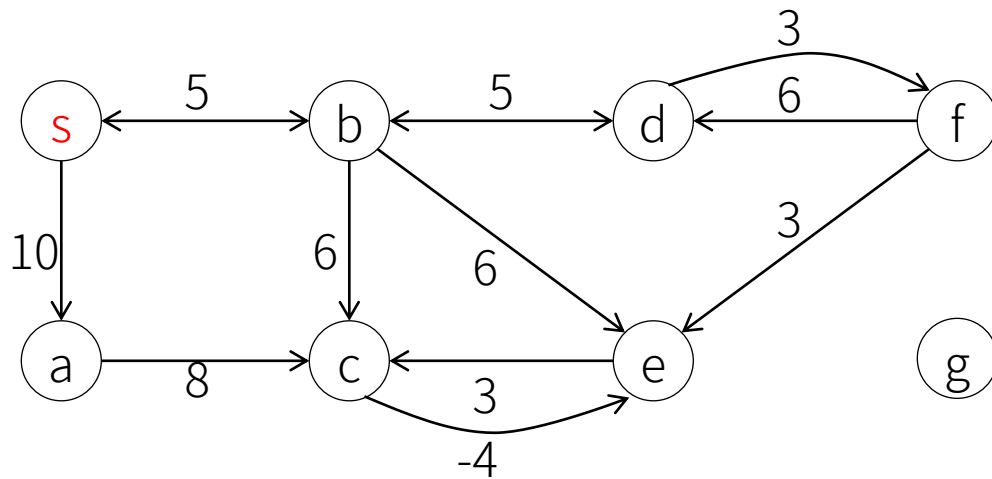
- Define $\delta(u, v) = \infty$ if v is unreachable from u
- Define $\delta(u, v) = -\infty$ if there exists a negative cycle on a path from u to v

Q: Is it correct that a shortest path has at most $|V| - 1$ edges?

Yes.

Having no cycle implies that a shortest path has at most $|V| - 1$ edges.

Practice



Destination v	Shortest path from s to v	Shortest path weight
a	s a	10
b		
c	NIL	$-\infty$
d		
e		
f	s b d f	13
g	NIL	∞

Variants of shortest-path problems

- **Single-source shortest-path problem:** Given a graph $G = (V, E)$ and a *source* vertex s in V , find the minimum cost paths from s to every vertex in V
- **Single-destination shortest-path problem:** Given a graph $G = (V, E)$ and a *destination* vertex t in V , find the minimum cost paths to t from every vertex in V
- **Single-pair shortest-path problem:** Find a shortest path from s to t for *given s and t*
- **All-pair shortest path problem:** Find a shortest path from s to t for *every pair of s and t*

Single-source shortest-path algorithms

- Dijkstra algorithm
 - Greedy
 - Requiring that all edge weights are **nonnegative**
- Bellman-Ford algorithm
 - Dynamic programming
 - General case, edge weights **may be negative**
- Both on a weighted, directed graph
- We'll introduce them next week

A very important technique: Relaxation

Many shortest-path algorithms work like this:

```
INITIALIZE-SINGLE-SOURCE (G, s)
  for v in G.V
    v.d =  $\infty$  //estimate
    v. $\pi$  = NIL //predecessor
  s.d = 0
```




Take a sequence of **relaxation** steps to update v.d and v. π



Output v.d and reconstruct shortest-paths from v. π

A very important technique: Relaxation


- The process of **relaxing an edge** (u, v)
= testing whether the shortest path weight of v **found so far** can be reduced by traveling over u
- 試試看經過 u 會不會比較好（更短的 $s - v$ 路徑）



Q: What if $w(u, v) = 10$?

A very important technique: Relaxation

- The process of **relaxing an edge** (u, v)
= testing whether the shortest path weight of v **found so far** can be reduced by traveling over u



```
RELAX( $u, v, w$ )  
  if  $v.d > u.d + w(u, v)$   
     $v.d = u.d + w(u, v)$   
     $v.\pi = u$ 
```


$v.d$ = shortest-path estimate

- An upper bound on $\delta(s, v)$ (Lemma 24.11)
- $v.d$ never increases during relaxation

$v.\pi$ = predecessor attribute

Recap: Optimal substructure

Shortest path problem (最短路徑問題) has optimal substructure



Path $P_{ac} + P_{cb}$ is a shortest path between A and B
 \Rightarrow Then P_{ac} must be a shortest path between A and C

Subpaths of shortest paths are shortest paths (Lemma 24.1)

Given a weighted, directed graph $G = (V, E)$ with weight function $w: E \rightarrow \mathbb{R}$, let $p = \langle v_0, v_1, \dots, v_k \rangle$ be a shortest path from vertex v_0 to vertex v_k and, for any i and j such that $0 \leq i \leq j \leq k$, let $p_{ij} = \langle v_i, v_{i+1}, \dots, v_j \rangle$ be the subpath of p from vertex i to vertex j . Then, p_{ij} is a shortest path from i to j .

Proof by contradiction

Triangle inequality (Lemma 24.10)

For any edge $(u, v) \in E$, $\delta(s, v) \leq \delta(s, u) + w(u, v)$

Upper-bound property (Lemma 24.11)

We always have $v.d \geq \delta(s, v)$ for all vertices $v \in V$, and once $v.d$ achieves the value $\delta(s, v)$, it never changes.

No-path property (Corollary 24.12)

If there is no path from s to v , then we always have $v.d = \delta(s, v) = \infty$

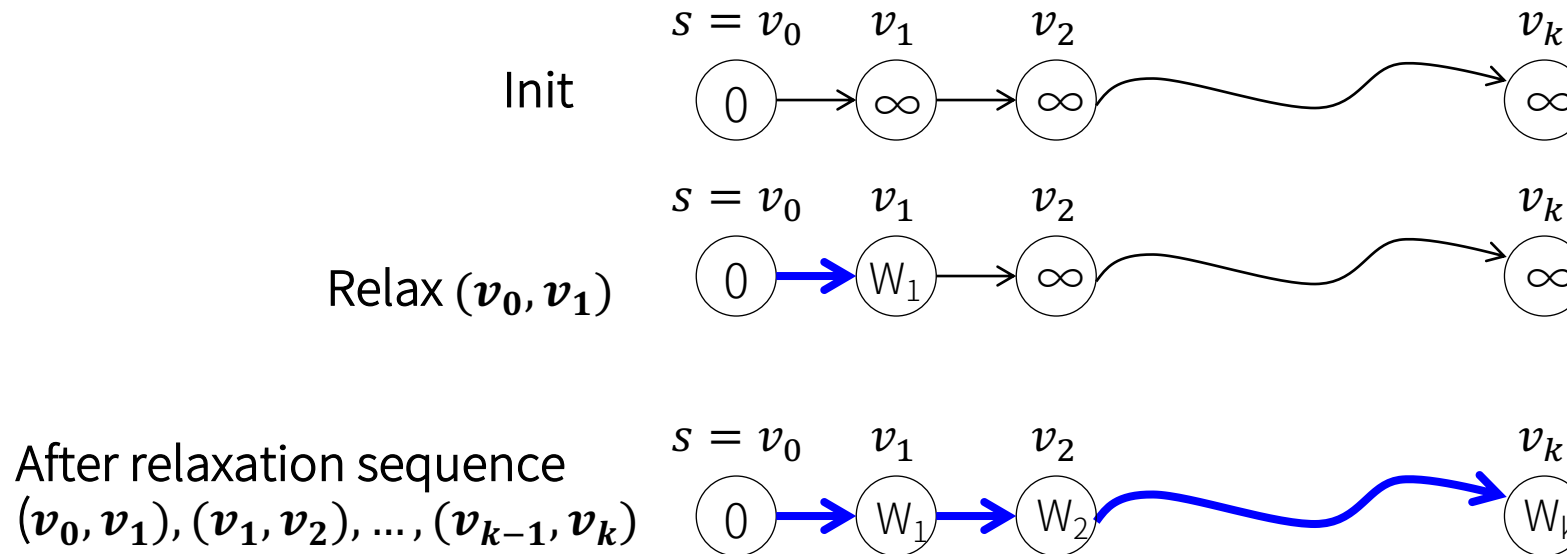
Convergence property (Lemma 24.14)

If $s \rightsquigarrow u \rightarrow v$ is a shortest path in G for some $u, v \in V$, and if $u.d = \delta(s, u)$ at any time prior to relaxing edge (u, v) , then $v.d = \delta(s, v)$ at all times afterward.


Path-relaxation property (Lemma 24.15)

$v_k.d = \delta(s, v_k)$ after relaxation sequence $(v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k)$

- Let $p = \langle v_0, v_1, \dots, v_k \rangle$ be a shortest path from $s = v_0$ to v_k
- Let $W_i = \sum_1^i (v_{i-1}, v_i)$, W_i be the shortest path weight $\delta(s, v_i)$ because of optimal substructure



Note: 此性質對於任何包含這個最短路徑邊的 relaxation sequence 都成立, e.g.,
 $(v_0, v_1), (a, b), (d, c), (v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k)$



• $\delta(s, e) = 9$

• A shortest path from s to $e = \langle s, b, d, f, e \rangle$

Q: After relaxing (s, b) , (b, d) , (d, f) , (f, e) in order, what's the value of $e.d$?

9

Q: Will the value of $e.d$ remain the same after relaxing the edges in a different order, such as (s, b) , (d, f) , (b, d) , (f, e) ?

Not necessary

Q: How about relaxing (s, b) , (b, e) , (s, a) , (b, d) , (d, f) , (e, c) , (f, e) ?

9

Predecessor-subgraph property (Lemma 24.17)

Once $v.d = \delta(s, v)$ for all $v \in V$, the predecessor subgraph is a shortest-paths tree rooted at s .

Shortest-paths tree

A **shortest-paths tree** $G' = (V', E')$ of s is a subgraph of G s.t.:

- V' is the set of vertices reachable from s in G
- G' forms a rooted tree with root s
- For all v in V' , the unique simple path from s to v in G' is a shortest path from s to v in G

